ObliVM: A Programming Framework for Secure Computation

Chang Liu*, Xiao Shaun Wang*, Kartik Nayak*, Yan Huang† and Elaine Shi*

*University of Maryland and †Indiana University

{liuchang,wangxiao,kartik,elaine}@cs.umd.edu, yh33@indiana.edu

(Conference version. Full version will appear shortly.)

Abstract—We design and develop **ObliVM**, a programming framework for secure computation. **ObliVM** offers a domain-specific language designed for compilation of programs into efficient oblivious representations suitable for secure computation. **ObliVM** offers a powerful, expressive programming language and user-friendly oblivious programming abstractions. We develop various showcase applications such as data mining, streaming algorithms, graph algorithms, genomic data analysis, and data structures, and demonstrate the scalability of **ObliVM** to bigger data sizes. We also show how **ObliVM** significantly reduces development effort while retaining competitive performance for a wide range of applications in comparison with hand-crafted solutions. We are in the process of open-sourcing **ObliVM** and our rich libraries to the community (www.oblivm.com), offering a reusable framework to implement and distribute new cryptographic algorithms.

I. INTRODUCTION

Secure computation [1], [2] is a powerful cryptographic primitive that allows multiple parties to perform rich data analytics over their private data, while preserving each individual or organization's privacy. The past decade has witnessed enormous progress in the practical efficiency of secure computation protocols [3]–[8] As a result, secure computation has evolved from being just a nice theoretical concept to having real system prototypes [9]–[17]. Several attempts to commercialize secure computation techniques have also been made [18], [19].

Architecting a system framework for secure computation presents numerous challenges. First, the system must allow *non-specialist programmers* without security expertise to develop applications. Second, *efficiency* is a first-class concern in the design space, and scalability to big data is essential in many interesting real-life applications. Third, the framework must be reusable: *expert programmers* should be able to easily extend the system with rich, optimized libraries or customized cryptographic protocols, and make them available to non-specialist application developers.

We design and build **ObliVM**, a system framework for automated secure multi-party computation. **ObliVM** is designed to allow non-specialist programmers to write programs much as they do today, and our ObliVM compiler compiles the program to an efficient secure computation protocol. To this end, ObliVM offers a domain-specific language that is intended to address a fundamental representation gap, namely, secure computation protocols (and other branches of modern cryptography) rely on *circuits* as an abstraction of computation, whereas real-life developers write *programs* instead. In architecting **ObliVM**, our main contribution is the design of programming support and compiler techniques that facilitate such program-to-circuit conversion while ensuring maximal efficiency. Presently, our framework assumes a semi-honest two-party protocol in the back end. To demonstrate an end-to-end system, we chose to implement an improved Garbled Circuit protocol as the back end, since it is among the most practical protocols to date. Our **ObliVM** framework, including source code and demo applications, will be open-sourced on our project website http://www.oblivm.com.

A. Background: "Oblivious" Programs and Circuits

To aid understanding, it helps to first think about an intuitive but somewhat imprecise view: Each variable and each memory location is labeled either as *secret* or *public*. Any secret variable or memory contents are secret-shared among the two parties such that neither party sees the values. The two parties run a cryptographic protocol to securely evaluate each instruction, making accesses to memory (public or secret-shared) whenever necessary. All messages transmitted are naturally secured by the underlying cryptographic protocol. However, the parties can additionally observe the following execution traces during the protocol execution: 1) the program counter (also referred to as the instruction trace); 2) addresses of all memory accesses (also referred to as the memory trace); and 3) the value of every public or declassified variable (similar to the notion of a low or declassified variable in standard information flow terminology). Imprecisely speaking, for security, it is imperative that the program's observable execution traces (not including the outcome) be "oblivious" to the secret inputs. A more formal security definition involves the use of a simulation paradigm that is standard in the cryptography literature [20], and is similar to the notion adopted in the SCVM work [15].

Relationship between oblivious programs and circuits. If a program is trace-oblivious by the aforementioned informal definition, it is then easy to convert the program into a sequence of circuits. These circuits are allowed to take memory accesses as inputs, however, these memory access must be oblivious to preserve security. By contrast, if a program is not memory-trace oblivious, then a dynamic memory access (whose address depends on secret inputs) cannot be efficiently made in the circuit representation – a straightforward approach (which is implicitly taken by almost all previous works except SCVM [15]) is to translate each dynamic memory access into a linear scan of memory in the resulting circuit, incurring prohibitive costs for large data sizes.

Moreover, instruction-trace obliviousness is effectively

guaranteed by executing both branches of a secret conditional in the resulting circuit where only one branch's execution takes effect. Our type system (formally defined in a separate manuscript [21]) rejects programs that loop on secret variables – in these cases, a maximum public bound on the loop guard can be supplied instead.

B. ObliVM Overview and Contributions

In designing and building ObliVM, we make the following contributions.

Programming abstractions for oblivious algorithms. The most challenging part about ensuring a program's obliviousness is memory-trace obliviousness – therefore our discussions below will focus on memory-trace obliviousness. A straightforward approach (henceforth referred to as the generic ORAM baseline) is to provide an Oblivious RAM (ORAM) abstraction, and require that all arrays (whose access patterns depend on secret inputs) be stored and accessed via ORAM. This approach, which was effectively taken by SCVM [15], is generic, but does not necessarily yield the most efficient oblivious implementation for each specific program.

At the other end of the spectrum, a line of research has focused on customized oblivious algorithms for special tasks (sometimes also referred to as circuit structure design). For example, efficient oblivious algorithms have been demonstrated for graph algorithms [22], [23], machine learning algorithms [24], [25], and data structures [26]–[28]. The customized approach can outperform generic ORAM, but is extremely costly in terms of the amount of cryptographic expertise and time consumed.

ObliVM aims to achieve the best of both worlds by offering oblivious programming abstractions that are both user- and compiler friendly. These programming abstractions are high-level programming constructs that can be understood and employed by non-specialist programmers without security expertise. Behind the scenes, ObliVM translates programs written in these abstractions into efficient oblivious algorithms that outperform generic ORAM. When oblivious programming abstractions are not applicable, ObliVM falls back to employing ORAM to translate programs to efficient circuit representations. Presently, ObliVM offers the following oblivious programming abstractions: MapReduce abstractions, abstractions for oblivious data structures, and a new loop coalescing abstraction which enables novel oblivious graph algorithms. We remark that this is by no means an exhaustive list of possible programming abstractions that facilitate obliviousness. It would be exciting future research to uncover new oblivious programming abstractions and incorporate them into our ObliVM framework.

An expressive programming language. ObliVM offers an expressive and versatile programming language called ObliVM-lang. When designing ObliVM-lang, we have the following goals.

- Non-specialist application developers find the language intuitive.
- Expert programmers should be able to extend our framework with new features. For example, an expert programmer should be able to introduce new, user-facing oblivious

programming abstractions by embedding them as libraries in ObliVM-lang (see Section IV-B for an example).
- Expert programmers can implement even low-level circuit libraries directly atop ObliVM-lang. Recall that unlike a programming language in the traditional sense, here the underlying cryptography fundamentally speaks only of AND and XOR gates. Even basic instructions such as addition, multiplication, and ORAM accesses must be developed from scratch by an expert programmer. In most previous frameworks, circuit libraries for these basic operations are developed in the back end. ObliVM, for the first time, allows the development of such circuit libraries in the source language, greatly reducing programming complexity. Section V-A demonstrates case studies for implementing basic arithmetic operations and Circuit ORAM atop our source language ObliVM.
- Expert programmers can implement customized protocols in the back end (e.g., faster protocols for performing big integer operations or matrix operations), and export these customized protocols to the source language as native types and native functions.

To simultaneously realize these aforementioned goals, we need a much more powerful and expressive programming language than any existing language for secure computation [10], [14]–[17]. Our ObliVM-lang extends the SCVM language by Liu et al. [15] and offers new features such as phantom functions, generic constants, random types, as well as native types and functions. We will show why these language features are critical for implementing oblivious programming abstractions and low-level circuit libraries.

Additional architectural choices. ObliVM also allows expert programmers to develop customized cryptographic protocols (not necessarily based on Garbled Circuit) in the back end. These customized back end protocols can be exposed to the source language through native types and native function calls, making them immediately reusable by others. Section VI describes an example where an expert programmer designs a customized protocol for `BigInteger` operations using additively-homomorphic encryption. The resulting `BigInteger` types and operations can then be exported into our source language ObliVM-lang.

C. Applications and Evaluation

ObliVM's easy programmability allowed us to develop a suite of libraries and applications, including streaming algorithms, data structures, machine learning algorithms, and graph algorithms. These libraries and applications will be shipped with the ObliVM framework. Our application-driven evaluation suggests the following results:

Efficiency. We use ObliVM's user-facing programming abstractions to develop a suite of applications. We show that over a variety of benchmarking applications, the resulting circuits generated by ObliVM can be orders of magnitude smaller than the generic ORAM baseline (assuming that the state-of-the-art Circuit ORAM [29] is adopted for the baseline) under moderately large data sizes. We also compare our ObliVM-generated circuits with hand-crafted designs, and show that for a variety of applications, our auto-generated circuits are only **0.5%** to **2%** bigger in size than oblivious algorithms hand-crafted by human experts.

Development effort. We give case studies to show how ObliVM greatly reduces the development effort and expertise needed to create applications over secure computation.

New oblivious algorithms. We describe a few new oblivious algorithms that we uncover during this process of programming language and algorithms co-design. Specifically, we demonstrate new oblivious graph algorithms including oblivious Depth-First-Search for dense graphs, oblivious shortest path for sparse graphs, and an oblivious minimum spanning tree algorithm.

D. Threat Model, Deployment, and Scope

Deployment scenarios and threat model. As mentioned, ObliVM presently supports a two-party semi-honest protocol. We consider the following primary deployment scenarios:

1) Two parties, Alice and Bob, each comes with their own private data, and engage in a two-party protocol. For example, Goldman Sachs and Bridgewater would like to perform joint computation over their private market research data to learn market trends.

2) One or more users break their private data (e.g., genomic data) into secret shares, and split the shares among two non-colluding cloud providers. The shares at each cloud provider are completely random and reveal no information. To perform computation over the secret-shared data, the two cloud providers engage in a secure 2-party computation protocol.

3) Similar as the above, but the two servers are within the same cloud or under the same administration. This can serve to mitigate Advanced Persistent Threats or insider threats, since compromise of a single machine will no longer lead to the breach of private data. Similar architectures have been explored in commercial products such as RSA's distributed credential protection [30].

In the first scenario, Alice and Bob should not learn anything about each other's data besides the outcome of the computation. In the second and third scenarios, the two servers should learn nothing about the users' data other than the outcome of the computation – note that the outcome of the computation can also be easily hidden simply by XORing the outcome with a secret random mask (like a one-time pad). We assume that the program text (i.e., code) is public.

Scope. A subset of ObliVM's source language ObliVM-lang has a security type system which, roughly speaking, ensures that the program's execution traces are independent of secret inputs [15], [31]. However, a formal treatment of the language and the type system is outside the scope of this paper and deferred to a forthcoming manuscript [21].

By designing a new language, ObliVM does not directly retrofit legacy code. Such a design choice maximizes opportunities for compile-time optimizations. We note, however, that in subsequent work joint with our collaborators [32], we have implemented a MIPS CPU in ObliVM, which can securely evaluate standard MIPS instructions in a way that leaks only the termination channel (i.e., total runtime of the program) – this secure MIPS CPU essentially provides backward compatibility atop ObliVM whenever needed.

II. RELATED WORK

Existing general-purpose secure computation systems can be classified roughly based on two mostly orthogonal dimensions: 1) which "back end" secure computation protocol they adopt – this will also decide whether the system is secure against semi-honest or malicious adversaries, and whether the system supports two or multiple parties; and 2) whether they offer programming and compiler support – and if so, which language and compiler they adopt.

A. Back End: Secure Computation Implementations

Below we discuss choices of back end secure computation protocols and implementations. As discussed later, under realistic bandwidth provisioning about 1.4MB/sec, Garbled Circuit is presently among the fastest general-purpose protocol for secure computation. Currently, ObliVM primarily supports a semi-honest Garbled Circuit based back end, but developers can introduce customized gadgets for special-purpose types and functions (e.g., operations on sets, matrices, and big integers), and export them as native types and functions in the source language. It would not be too hard to extend ObliVM to support additional back end protocols such as GMW and FHE – in particular, almost all known protocols use a circuit abstraction (either boolean or arithmetic circuits). An interesting direction of the future research is to create new, compile-time optimizations that automatically selects the optimal mix of protocols for a given program, similar to what TASTY [13] proposed, but in a much broader sense.

Garbled Circuit (GC) implementations. The Garble Circuit protocol was first proposed by Andrew Yao [35]. Numerous later works improved the original protocol: Free XOR shows that XOR gates can be computed almost "freely" [5]–[7]. Row reduction techniques show that only 2 or 3 garbled entries (rather than 4) need to be sent across the network per AND gate [36], [37]. A building block called Oblivious Transfer (OT) that is necessary for Garbled Circuit protocols was proposed and improved in a sequence of works as well [3], [8].

Several works have implemented the Garbled Circuit protocol – we give an overview of their features and performance characteristics in Table I.

Non-GC protocols and implementations. Besides Garbled Circuits, several other techniques have been proposed for general-purpose secure computation, including FHE [38], GMW [2], schemes based on linear secret-sharing [9], [14], etc. More discussions on non-GC protocols and implementations can be found in our online technical report [39].

B. Programming and Compiler Support

Secure computation compilers are in charge of compiling programs to circuit representations. One subtlety must be clarified: instead of a single circuit, here a program may be compiled to a sequence of circuits whose inputs are oblivious memory accesses. The number of these circuits will determine the number of interactions of the protocol.

Circuit generation. One key question is whether the circuits are *fully materialized* or *generated on the fly* during secure

GC Back End	Features	Garbling Speed	Bandwidth to match compute	Adopted by
Fairplay [12]	Java-based	$\leq$ 30 gates/sec	900Bps	
FastGC [33]	Java-based	96K gates/sec	2.8MBps	CBMC-GC [16] PCF [10] SCVM [15]
ObliVM-GC (this paper)	Java-based	670K gates/sec, 1.8M gates/sec (online)	19.6MBps 54MBps (online)	ObliVM GraphSC [24]
GraphSC [24] (extends ObliVM-GC)	Java-based Parallelizable	580K gates/sec per pair of cores 1.4M gates/sec per pair of cores (online)	16MBps per pair of cores 41MBps per pair of cores (online)	
JustGarble [4]	C-based Hardware AES-NI Garbling only, does not run end-to-end	11M gates/sec	315MBps	TinyGarble [34]
KSS [11]	Parallel execution in malicious mode Hardware AES-NI	320 gates/sec per pair of cores	2.4MBps per pair of cores	PCF [10]

TABLE I: **Summary of known (2-party) Garbled Circuit back ends.** The gates/sec metric refer specifically to AND gates, since XOR gates are considered free [5]–[7]. Measurements for different papers are taken on off-the-shelf computers representative of when each paper is written. ObliVM essentially adopts a much better architected and engineered version of FastGC [33]. The focus of this paper is our language, programming abstractions, and compiler. It is our future work to extend JustGarle (C-based, hardware AES-NI) to a fully working back end and integrate it with our language and compiler.

computation. Many first-generation secure computation compilers such as Fairplay [12], TASTY [13], Sharemind [9], CBMC-GC [16], PICCO [14], KSS12 [11] generate target code containing the fully materialized circuits. This approach has the following drawbacks. First, the target code size and compile time are proportional to the circuit size. That is why some works report large compile times (e.g., 8.2 seconds for a circuit of size 700K in KSS12 [11]). Second, the program must be recompiled for every input data size – possibly taking a long time again!

Newer generations of secure computation compilers (e.g., PCF, Wysteria, and SCVM [10], [15], [17]) employ *program-style target code* instead. Program-style target code is essentially a more compact intermediate representation of circuits – fundamentally, the succinctness comes from introducing looping instructions in the intermediate representation, such that the circuit need not be fully unrolled in this intermediate representation.

The resulting program-style target code can then be securely evaluated using a cryptographic protocol such as Garbled Circuit or GMW. Typically these protocols perform per-gate computation – therefore, circuits are effectively generated on-the-fly at runtime. ObliVM also adopts program-style target code and on-the-fly circuit generation. Specifically, the circuit generation is pipelined using a well-known technique by Huang et al. [33] such that the circuit is never materialized entirely, and thus only a constant amount of working memory is necessary. Further, we stress that on-the-fly circuit generation incurs unnoticeable cost in comparison with the time required to compute the cryptographic protocol. In ObliVM, on-the-fly circuit generation only contributes to less than 0.1% of the total runtime.

Finally, in a concurrent work called TinyGarble [34], Songhori et al. show that by partially materializing a circuit, they can have a somewhat more global view of the circuit. Thus they show how to borrow hardware circuit synthesis

techniques to optimize the circuit size by roughly 50% to 80% in comparison with PCF [10]. TinyGarble's techniques are orthogonal and complementary to this work.

ORAM support. Almost all existing secure computation compilers, including most recent ones such as Wysteria [17], PCF [10], and TinyGarble [34], compile dynamic memory accesses (whose addresses depend on secret inputs) to a linear scan of memory in the circuit representation. This is completely unscalable for big data sizes. A solution to this problem lies in Oblivious RAM (ORAM), first proposed by Goldreich and Ostrovsky [40], [41]. To the best of our knowledge, the only known compiler that provides ORAM support is our prior work SCVM which ObliVM builds on. SCVM employs the binary-tree ORAM [42] to implement dynamic memory accesses. Presently, Circuit ORAM is the most efficient ORAM scheme for secure computation – and ObliVM is the first to offer a Circuit ORAM implementation.

Language expressiveness and formal security. Most existing languages for secure computation are restrictive in nature. Existing languages [10], [13]–[17] lack essential features such as function calls and public loops inside secret-`if`s. This prevents the implementation of a large class of interesting programs. We also offer several other new features such as native primitives, random types (with an affine type system), and generic constants that were lacking in previous languages [10], [13]–[17].

Earlier domain-specific languages [10], [13], [14], [16] for secure computation do not aim to offer formal security. More recent languages such as SCVM [15] and Wysteria [17] offer formal security through new type systems. In comparison, Wysteria's type system is too restrictive – for example, Wysteria rejects programs with public loops and function calls inside secret-`if`s. This prevents many interesting applications – for example, it is not feasible to implement ORAM and oblivious data structures efficiently in Wysteria. On the other hand, Wysteria supports multiple parties, and abstractions for writing

code generic in number of parties. Extending to multiple parties is future work for ObliVM.

III. PROGRAMMING LANGUAGE AND COMPILER

As mentioned earlier, we wish to design a powerful source language ObliVM-lang such that an expert programmer can *i)* develop oblivious programming abstractions as libraries and offer them to non-specialist programmers; and *ii)* implement low-level circuit gadgets atop ObliVM-lang.

ObliVM-lang builds on top of the recent SCVM source language [15] – the only known language to date that supports ORAM abstractions, and therefore offers scalability to big data. In this section, we will describe new features that ObliVM-lang offers and explain intuitions behind our security type system which is formalized in a separate manuscript [21].

As compelling applications of ObliVM-lang, in Section IV, we give concrete case studies and show how to implement oblivious programming abstractions and low-level circuit libraries atop ObliVM-lang.

A. Language features for expressiveness and efficiency

Security labels. Except for the new random type introduced in Section III-B, all other variables and arrays are either of a `public` or `secure` type. `secure` variables are secret-shared between the two parties such that neither party sees the value. `public` variables are observable by both parties. Arrays can be publicly or secretly indexable. For example,

- `secure int10[public 1000] keys`: secret array contents but indices to the array must be public. This array will be secret shared but not placed in ORAMs.
- `secure int10[secure 1000] keys`: This array will be placed in a secret-shared ORAM, and we allow secret indices into the array.

Standard features. ObliVM-lang allows programmers to use C-style keyword `struct` to define *record types*. It also supports *generic types* similar to templates in C++. For example, a binary tree with public topological structure but secret per-node data can be defined without using pointers (assuming its capacity is 1000 nodes):

```
struct KeyValueTable<T> {
  secure int10[public 1000] keys;
  T[public 1000] values;
};
```

In the above, the type `int10` means that its value is a 10-bit signed integer. Each element in the array `values` has a generic type `T` similar to C++ templates. ObliVM-lang assumes data of type `T` to be secret-shared. In the future, we will improve the compiler to support public generic types.

Generic constants. Besides general types, ObliVM-lang also supports *generic constants* to further improve the reusability. Let us consider the following tree example:

```
struct TreeNode@m<T> {
  public int@m key;
  T value;
  public int@m left, right;
};
struct Tree@m<T> {
  TreeNode<T>[public (1<<m)-1] nodes;
  public int@m root;
};
```

This code defines a binary search tree implementation of a key-value store, where keys are m-bit integers. The *generic constant* `@m` is a variable whose value will be instantiated to a constant. It hints that `m` bits are enough to represent all the position references to the array. The type `int@m` refers to an integer type with `m` bits. Further, the capacity of array `nodes` can be determined by `m` as well (i.e. `(1<<m)-1`). Note that Zhang et al. [14] also allow specifying the length of an integer, but require this length to be a hard-coded constant – this necessitates modification and recompilation of the program for different inputs. ObliVM-lang's generic constant approach eliminates this constraint, and thus improves reusability.

Functions. ObliVM-lang allows programmers to define functions. For example, following the `Tree` defined as above, programmers can write a function to search the value associated with a given key in the tree as follows:

```
1   T Tree@m<T>.search(public int@m key) {
2     public int@m now = this.root, tk;
3     T ret;
4     while (now != -1) {
5       tk = this.nodes[now].key;
6       if (tk == key)
7         ret = this.nodes[now].value;
8       if (tk <= key)
9         now = this.nodes[now].right;
10      else
11        now = this.nodes[now].left;
12    }
13    return ret
14  };
```

This function is a method of a `Tree` object, and takes a `key` as input, and returns a value of type `T`. The function body defines three local variables `now` and `tk` of type `public int@m`, and `ret` of type `T`. The definition of a local variable (e.g. `now`) can be accompanied with an optional initialization expression (e.g. `this.root`). When a variable (e.g. `ret` or `tk`) is not initialized explicitly, it is initialized to be a default value depending on its type.

The rest of the function is standard, C-like code, except that ObliVM-lang requires exactly one return statement at the bottom of a function whose return type is not `void`. We highlight that ObliVM-lang allows arbitrary looping on a public guard (e.g. line 4) without loop unrolling, which cannot be compiled in previous loop-elimination-based work [9], [11]–[14], [16].

Function types. Programmers can define a variable to have function type, similar to function pointers in C. To avoid the complexity of handling arbitrary higher order functions, the input and return types of a function type must not be function types. Further, generic types cannot be instantiated with function types.

Native primitives. ObliVM-lang supports native types and native functions. For example, ObliVM-lang's default back end implementation is ObliVM-GC, which is implemented in Java. Suppose an alternative BigInteger implementation in ObliVM-GC (e.g., using additively homomorphic encryption) is available in a Java class called `BigInteger`. Programmers can define

```
typedef BigInt@m = native BigInteger;
```

Suppose that this class supports four operations: `add`, `multiply`, `fromInt` and `toInt`, where the first two operations are arithmetic operations and last two operations are used to convert between Garbled Circuit-based integers and HE-based integers. We can expose these to the source language by declaring:

```
BigInt@m BigInt@m.add(BigInt@m x,
    BigInt@m y)= native BigInteger.add;
BigInt@m BigInt@m.multiply(BigInt@m x,
   = BigInt@m y) native BigInteger.multiply;
BigInt@m BigInt@m.fromInt(int@m y)
   = native BigInteger.fromInt;
int@m BigInt@m.toInt(BigInt@m y)
   = native BigInteger.toInt;
```

B. Language features for security

The key requirement of ObliVM-lang is that a program's execution traces will not leak information. These execution traces include a memory trace, an instruction trace, a function stack trace, and a declassification trace. The trace definitions are similar to Liu et al. [15]. We develop a security type system for ObliVM-lang.

Liu et al. [15] has discussed how to prevent memory traces and instruction traces from leaking information. We explain the basic ideas of ObliVM-lang's type system concerning functions and declassifications, but defer a formal discussion to a separate manuscript [21].

Random numbers and implicit declassifications. Many oblivious programs such as ORAM and oblivious data structures crucially rely on randomness. In particular, their obliviousness guarantee has the following nature: the joint *distribution* of memory traces is identical regardless of secret inputs (these algorithms typically have a cryptographically negligible probability of correctness failure). ObliVM-lang supports reasoning of such "distributional" trace-obliviousness by providing *random types* associated with an affine type system. For instance, `rnd32` is the type of a 32-bit random integer. A random number will always be secret-shared between the two parties.

To generate a random number, there is a built-in function `RND` with the following signature:

```
rnd@m RND(public int32 m)
```

This function takes a public 32-bit integer `m` as input, and returns `m` random bits. Note that `rnd@m` is a *dependent type*, whose type depends on values, i.e. `m`. To avoid the complexity of handling general dependent types, the ObliVM-lang compiler restricts the usage of dependent types to only this built-in function, and handles it specially.

In our ObliVM framework, outputs of a computation can be explicitly declassified with special syntax. Random numbers

are allowed *implicit declassification* – by assigning them to public variables. Here "implicitness" means that the declassification happens not because this is a specified outcome of the computation.

For security, we must ensure that each random number is implicitly declassified *at most once* for the following reason. When implicitly declassifying a random number, both parties observe the random number as part of the trace. Now consider the following example where `s` is a secret variable.

```
1    rnd32 r1 = RND(32), r2= RND(32);
2    public int32 z;
3    if (s) z = r1; // implicit declass
4    else z = r2; // implicit declass
     ......
XX   public int32 y = r2; // NOT OK
```

In this program, random variables `r1` and `r2` are initialized in Line 1 – these variables are assigned a fresh, random value upon initialization. Up to Line 4 , random variables `r1` and `r2` are each declassified no more than once. Line XX, however, could potentially cause `r2` to be declassified more than once. Line XX clearly is not secure since in this case the observable public variable `y` and `z` could be correlated – depending on which secret branch was taken earlier.

Therefore, we use an *affine type* system to ensure that each random variable is implicitly declassified at most once. This way, each time a random variable is implicitly declassified, it will introduce a independently uniform variable to the observable trace. In our security proof, a simulator can just sample this random number to simulate the trace.

It turns out that the above example reflects the essence of what is needed to implement oblivious RAM and oblivious data structures in our source language. We refer the readers to Sections IV and V-B for details.

Function calls and phantom functions. A straightforward idea to prevent stack behavior from leaking information is to enforce function calls in a public context. Then the requirement is that each function's body must satisfy memory- and instruction-trace obliviousness. Further, by defining native functions, ObliVM-lang implicitly assumes that their implementations satisfy memory- and instruction-trace obliviousness.

Beyond this basic idea, ObliVM-lang makes a step forward to enabling *function calls within a secret if-statement* by introducing the notion of *phantom function*. The idea is that each function can be executed in dual modes, a *real* mode and and a *phantom* mode. In the real mode, all statements are executed normal with real computation and real side effects. In the phantom mode, the function execution merely simulates the memory traces of the real world; no side effects take place; and the phantom function call returns a secret-shared default value of the specified return type. This is similar to padding ideas used in several previous works [43], [44].

We will illustrate the use of phantom function with the following `prefixSum` example. The function `prefixSum`(n) accesses a global integer array `a`, and computes the prefix sum of the first $n + 1$ elements in `a`. After accessing each element (Line 3), the element in array `a` will be set to 0 (Line 4).

```
1   phantom secure int32 prefixSum
2       (public int32 n) {
3     secure int32 ret=a[n];
4     a[n]=0;
5     if (n != 0) ret = ret+prefixSum(n-1);
6     return ret;
7   }
```

The keyword `phantom` indicates that the function `prefixSum` is a phantom function.

Consider the following code to call the phantom functions:

```
if (s) then x = prefixSum(n);
```

To ensure security, `prefixSum` will always be called no matter s is true or false. When s is false, however, it must be guaranteed that (1) elements in array a will not be assigned to be 0; and (2) the function generates traces with the same probability as when s is true. To this end, the compiler will generate target code with the following signature:

```
prefixSum(idx, indicator)
```

where `indicator` means whether the function will be called in the real or phantom mode. To achieve the first goal, the global variable will be modified only if `indicator` is false. The compiler will compile the code in line 4 into the following pseudo-code:

```
a[idx]=mux(0, a[idx], indicator);
```

It is easy to see, that all instructions will be executed, and thus the generated traces are identical regardless of the value of `indicator`. Note, that such a function is not implementable in any prior loop-unrolling based compiler, since n is provided at runtime only.

IV. User-Facing Oblivious Programming Abstractions

Programming abstractions such as MapReduce and GraphLab have been popularized in the parallel computing domain. In particular, programs written for a traditional sequential programming paradigm are difficult to parallelize automatically by an optimizing compiler. These new paradigms are not only easy for users to understand and program with, but also provide insights on the structure of the problem, and facilitate parallelization in an automated manner.

In this section, we would like to take a similar approach towards oblivious programming as well. The idea is to develop oblivious programming abstractions that can be easily understood and consumed by non-specialist programmers, and our compiler can compile programs into efficient oblivious algorithms. In comparison, if these programs were written in a traditional imperative-style programming language like C, compile-time optimizations would have been much more limited.

A. MapReduce Programming Abstractions

An interesting observation is that "parallelism facilitates obliviousness" [45], [46]. If a program (or part of a program) can be efficiently expressed in parallel programming paradigms such as MapReduce and GraphLab [47], [48] (with a few additional constraints), there is an efficient oblivious algorithm to compute this task. We stress that in this paper, we consider MapReduce merely as a programming abstraction that facilitates obliviousness – in reality we compile MapReduce programs to *sequential* implementations that runs on a single thread. Parallelizing the algorithms is outside the scope of this paper. However, in a subsequent work GraphSC [24] jointly with our collaborators, we do offer parallel oblivious implementations of programs written in a GraphLab abstraction – and doing so requires the design of new, non-trivial *parallel oblivious algorithms* detailed in the GraphSC paper [24].

Background: Oblivious algorithms for streaming MapReduce. A *streaming* MapReduce program consists of two basic operations, `map` and `reduce`.

- The `map` operation: takes an array denoted $\{\alpha_i\}_{i \in [n]}$ where each $\alpha_i \in \mathcal{D}$ for some domain $\mathcal{D}$, and a function `mapper` : $\mathcal{D} \to \mathcal{K} \times \mathcal{V}$. Now `map` would apply $(k_i, v_i) := \mathtt{mapper}(\alpha_i)$ to each α_i, and output an array of key-value pairs $\{(k_i, v_i)\}_{i \in [n]}$.
- The `reduce` operation: takes in an array of key-value pairs denoted $\{(k_i, v_i)\}_{i \in [n]}$ and a function `reducer` : $\mathcal{K} \times \mathcal{V}^2 \to \mathcal{V}$. For every unique key k value in this array, let $(k, v_{i_1}), (k, v_{i_2}), \dots (k, v_{i_m})$ denote all occurrences with the key k. Now the `reduce` operation applies the following operation in a streaming fashion:

$$R_k := \mathtt{reducer}(k, \dots \mathtt{reducer}(k, \mathtt{reducer}(k, v_{i_1}, v_{i_2}), v_{i_3}), \dots, v_{i_m})$$

The result of the `reduce` operation is an array consisting of a pair (k, R_k) for every unique k value in the input array.

Goodrich and Mitzenmacher [45] observe that any program written in a streaming MapReduce abstraction can be converted to efficient oblivious algorithms, and they leverage this observation to aid the construction of an ORAM scheme.

- The `map` operation is inherently oblivious, and can be done by making a linear scan over the input array.
- The `reduce` operation can be made oblivious through an oblivious sorting (denoted o-sort) primitive.
  - First, o-sort the input array in ascending order of the key, such that all pairs with the same key are grouped together.
  - Next, in a single linear scan, apply the `reducer` function: *i)* If this is the last key-value pair for some key k, write down the result of the aggregation (k, R_k). *ii)* Else, write down a dummy entry $\bot$.
  - Finally, o-sort all the resulting entries to move $\bot$ to the end.

Providing the streaming MapReduce abstraction in ObliVM. It is easy to implement the streaming MapReduce abstraction as a library in our source language ObliVM-lang. The ObliVM-lang implementation of streaming MapReduce paradigm is provided in Figure 1.

MapReduce has two generic constants, `m` and `n`, to represent the sizes of the input and output respectively. It also has three generic types, `I` for inputs' type, `K`, for output keys' type, and `V`, for output values' type. All of these three types are assumed to be secret.

```
 1   Pair<K,V>[public n] MapReduce@m@n<I,K,V>          18        } else {
 2      (I[public m] data,  Pair<K, V> map(I),        19          res[i].k.val = 0;
 3       V reduce(K, V, V), V initialVal,             20          key = d2[i+1].k;
 4       int2 cmp(K, K)) {                            21          val = initialVal;
 5     public int32 i;                                22        }
 6     Pair<K, V>[public m] d2;                       23        val = reduce(key, val, d2[i+1].v);
 7     for (i=0; i<m; i=i+1)                          24      }
 8       d2[i] = map(data[i]);                        25      res[m-1].k.val = 0;
 9     sort@m<K, V>(d2, 1, cmp);                      26      res[m-1].v.k = key;
10     K key = d2[0].k;                               27      res[m-1].v.v = val;
11     V val = initialVal;                            28      sort@m<int1, Pair<K, V>>
12     Pair<int1, Pair<K, V>>[public m] res;          29        (res, 1, zeroOneCmp);
13     for (i=0; i+1<m; i=i+1) {                      30      Pair<K, V>[public n] top;
14       res[i].v.k = key;                            31      for (i=0; i < n; i = i + 1)
15       res[i].v.v = val;                            32        top[i] = res[i].v;
16       if (cmp(key, d2[i+1].k)==0) {                33      return top;
17         res[i].k.val = 1;                          34   }
```

Fig. 1: **Streaming MapReduce in ObliVM-lang.** See Section IV-A for oblivious algorithms for the streaming MapReduce paradigm [45].

It takes five inputs, data for the input data, map for the mapper, reduce for the reducer, initialVal for the initial value for the reducer, and cmp to compare two keys of type K.

Lines 6-10 are the mapper phase of the algorithm, then line 11 uses the function sort to sort the intermediate results based on their keys. After line 11, the intermediate results with the same key are grouped together, and line 12-29 produce the output of the reduce phase with some dummy outputs. Finally, lines 30-35 use oblivious sort again to eliminate those dummy outputs, and eventually line 36 returns the final results.

Notice that in these functions, there are three arrays, data, d2, and res. The program declares all of them to have only public access pattern, because they are accessed by either a sequential scan, or an oblivious sorting. In this case, the compiler will not place these arrays into ORAM banks.

Using MapReduce. Figure 1 needs to be written by an expert developer only once. From then on, an end user can make use of this programming abstraction.

We further illustrate how to use the above MapReduce program to implement a histogram. In SCVM [15], a histogram program is as below.

```
for (public int i=0; i<n; ++i) c[i] = 0;
for (public int i=0; i<m; ++i) c[a[i]] ++;
```

This program counts the frequency of each values in $[0..n-1]$ in the array a of size m. Since the program makes dynamic memory accesses, the SCVM compiler would decide to put the array c inside an ORAM.

An end user can write the same program using a simple MapReduce abstraction as follows. Our ObliVM-lang compiler would generate target code that relies on oblivious sorting primitives rather than generic ORAM, improving the performance by a logarithmic factor in comparison with the SCVM implementation. In Section VII, we show that the practical performance gain ranges from $10\times$ to $400\times$.

```
int2 cmp(int32 x, int32 y) {
  int2 r = 0;
  if (x < y) r = -1;
  else if (x > y) r = 1;
  return r;
}
Pair<int32, int32> mapper(int32 x) {
  return Pair<int32, int32>(x, 1);
}
int32 reducer(int32 k, int32 v1, int32 v2) {
  return v1 + v2;
}
```

The top-level program can launch the computation using

```
c = MapReduce@m@n<int32, int32, int32>
  (a, mapper, reducer, cmp, 0);
```

B. Programming Abstractions for Data Structures

We now explain how to provide programming abstractions for a class of pointer-based oblivious data structures described by Wang et al. [26]. Figure 2b gives an example, where an expert programmer provides library support for implementing a class of pointer-based data structures such that a non-specialist programmer can implement data structures which will be compiled to efficient oblivious algorithms that outperform generic ORAM. We stress that while we give a stack example for simplicity, this paradigm is also applicable to other pointer-based data structures, such as AVL tree, heap, and queue.

Implementing oblivious data structure abstractions in ObliVM. We assume that the reader is familiar with the oblivious data structure algorithmic techniques described by Wang et al. [26]. To support efficient data structure implementations, an expert programmer implements two important objects (see Figure 2b):

- A Pointer object stores two important pieces of information: an index variable that stores the logical identifier of the memory block pointed to (each memory block has a globally unique index); and a pos variable that stores

```
1  struct StackNode@m<T> {
2    Pointer@m next;
3    T data;
4  };
5  struct Stack@m<T> {
6    Pointer@m top;
7    SecStore@m store;
8  };
9  phantom void Stack@m<T>.push(T data) {
10   StackNode@m<T> node = StackNode@m<T> (
       top, data);
11   this.top = this.store.allocate();
12   store.add(top.(index, pos), node);
13 }
14 phantom T Stack@m<T>.pop() {
15   StackNode@m<T> res = store
       .readAndRemove(top.(index, pos));
16   top = res.next;
17   return res.data;
18 }
```

(a) Oblivious stack by non-specialist programmers.

```
1  rnd@m RND(public int32 m) = native lib.rand;
2  struct Pointer@m {
3    int32 index;
4    rnd@m pos;
5  };
6  struct SecStore@m<T> {
7    CircuitORAM@m<T> oram;
8    int32 cnt;
9  };
10 phantom void SecStore@m<T>.add(int32 index,
     int@m pos, T data) {
11   oram.add(index, pos, data);
12 }
13 phantom T SecStore@m<T>
       .readAndRemove(int32 index, rnd@m pos) {
14   return oram.readAndRemove(index, pos);
15 }
16 phantom Pointer@m SecStore@m<T>.allocate() {
17   cnt = cnt + 1;
18   return Pointer@m(cnt, RND(m));
19 }
```

(b) Code by expert programmers to help non-specialists implement oblivious data structures.

Fig. 2: **Programming abstractions for oblivious data structures.**

the random leaf label in the ORAM tree of the memory block.

- A `SecStore` object essentially implements an ORAM, and provides the following member functions to an end-user: The `SecStore.remove` function essentially is a syntactic sugar for the ORAM's `readAndRemove` interface [29], [42], and the `SecStore.add` function is a syntactic sugar for the ORAM's `Add` interface [29], [42]. Finally, the `SecStore.allocate` function returns a new `Pointer` object to the caller. This new `Pointer` object is assigned a globally unique logical identifier (using a counter `cnt` that is incremented each time), and a fresh random chosen leaf label `RND`(m).

Stack implementation by a non-specialist programmer. Given abstractions provided by the expert programmer, a non-specialist programmer can now implement a class of data structures such as stack, queue, heap, AVL Tree, etc. Figure 2a gives a stack example.

Role of affine type system. We use Figure 2b as an example to illustrate how our `rnd` types with their affine type system can ensure security. As mentioned earlier, `rnd` types have an affine type system. This means that each `rnd` can be declassified (i.e., made public) at most once. In Figure 2b, the `oram.readAndRemove` call will declassify its argument `rnd@m pos` inside the implementation of the function body. From an algorithms perspective, this is because the leaf label `pos` will be revealed during the `readAndRemove` operation, incurring a memory trace where the value `rnd@m pos` will be observable by the adversary.

C. Loop Coalescing and New Oblivious Graph Algorithms

We introduce a new programming abstraction called loop coalescing, and show how this programming abstraction al-

lowed us to design novel oblivious graph algorithms such as Dijkstra's shortest path and minimum spanning tree for sparse graphs. Loop coalescing is non-trivial to embed as a library in ObliVM-lang. We therefore support this programming abstraction by introducing special syntax and modifications to our compiler. Specifically, we introduce a new syntax called *bounded-for* loop as shown in Figure 3. *For succinctness, in this section, we will present pseudo-code.*

In the program in Figure 3, the `bwhile`(n) and `bwhile`(m) syntax at Lines 1 and 3 indicate that the outer loop will be executed for a total of n times, whereas the inner loop will be executed for a total of m times – over all iterations of the outer loop.

To deal with loop coalescing, the compiler partitions the code within an bounded-loop into code blocks, each of which will branch at the end. The number of execution times for each code block will be computed as the bound number for the inner most bounded-loop that contains the code block. Then the compiler will transform a bounded loop into a normal loop, whose body simulates a state machine that each state contains a code block, and the branching statement at the end of each code block will be translated into an assignment statement that moves the state machine into a next state. The total number of iterations of the emitted normal loop is the summation of the execution times for all code blocks. Figure 3 illustrates this compilation process.

We now show how this loop coalescing technique leads to new novel oblivious graph algorithms.

Oblivious Dijkstra shortest path for sparse graphs. It is an open problem how to compute single source shortest path (SSSP) obliviously for *sparse graphs* more efficiently than generic ORAM approaches. Blanton et al. [49] designed a solution for a *dense* graph, but their technique cannot be

```
1 bwhile(n) (; u<n;) {              bwhile(n) (; u<n;) {          n + m + n        state = (u<n) ? 1 : -1;
2   total = total + 1;               total = total + 1;          iterations        for (__itr=0; __itr<n+m+n; __itr++) {
3   i=s[u];                          i=s[u];                     in total              if (state==1) {  total=total+1; i=s[u];
4   bwhile (m) (i<s[u+1]) {          bwhile (m) (i<s[u+1]) {    Block 1 ×n              state = (i<s[u+1]) ? 2 : 3;
5     // do something                  // do something                               } else if (state==2) { // do something
6     i=i+1;                           i=i+1;                                             i=i+1; state = (i<s[u+1]) ? 2 : 3;
7   }                                }                           Block 2 ×m          } else if (state==3) {
8   u=u+1;                           u=u+1;                                               u=u+1; state = (u<n) ? 1 : -1;
9 }                                }                             Block 3 ×n          } // else execution is finished
                                                                                   }
```
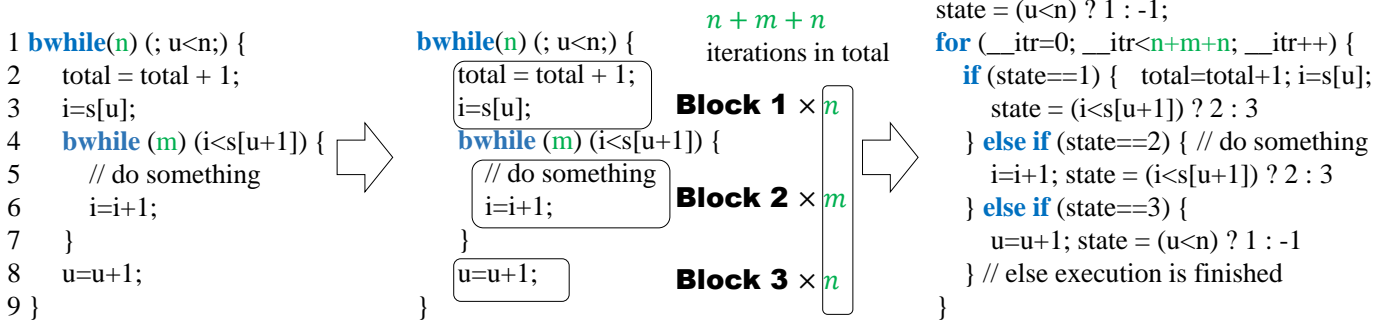
Fig. 3: **Loop coalescing.** The outer loop will be executed at most n times in total, the inner loop will be executed at most m times in total – over all iterations of the outer loop. A naive approach compiler would pad the outer and inner loop to n and m respectively, incurring $O(nm)$ cost. Our loop coalescing technique achieves $O(n + m)$ cost instead.

Algorithms		Complexity		
		Our Complexity	Generic ORAM	Best Known
Sparse Graph	Dijkstra's Algorithm	$O((E + V)\log^2 V)$	$O((E + V)\log^3 V)$	$O((E + V)\log^3 V)$ (Generic ORAM baseline [29])
	Prim's Algorithm	$O((E + V)\log^2 V)$	$O((E + V)\log^3 V)$	$O(E\frac{\log^3 V}{\log\log V})$ for $E = O(V\log^\gamma V), \gamma \geq 0$ [22]
				$O(E\frac{\log^3 V}{\log^\delta V})$ for $E = O(V2^{\log^\delta V}), \delta \in (0,1)$ [22]
				$O(E\log^2 V)$ for $E = \Omega(V^{1+\epsilon}), \epsilon \in (0,1]$ [22]
Dense Graph	Depth First Search	$O(V^2\log V)$	$O(V^2\log^2 V)$	$O(V^2\log^2 V)$ [49]

TABLE II: **Summary of algorithmic results.** All costs reported are in terms of circuit size. The asymptotic notation omits the bit-length of each word for simplicity. Our oblivious Dijkstra's algorithm and oblivious Prim's algorithm can be composed using our novel loop coalescing programming abstraction and oblivious data structures. Our oblivious DFS algorithm requires independent novel techniques. Due to space constraint, we only describe the oblivious Dijkstra's algorithm as an example of loop coalescing. We defer the full description of our oblivious MST and DFS algorithms to Appendix A.

Algorithm 1 Dijkstra' algorithm with bounded for

Secret Input: s: the source node
Secret Input: e: concatenation of adjacency lists stored in a single ORAM array. Each vertex's neighbors are stored adjacent to each other.
Secret Input: s[u]: sum of out-degree over vertices from 1 to u.
Output: dis: the shortest distance from source to each node
1: dis := $[\infty, \infty, ..., \infty]$
2: PQ.push(0, s)
3: dis[s] := 0
4: **bwhile**(V)(!PQ.empty())
5:     (dist, u) := PQ.deleteMin()
6:     **if**(dis[u] == dist) **then**
7:         dis[u] := −dis[u];
8:         **bfor**(E)(i := s[u]; i < s[u + 1]; i = i + 1)
9:             (u, v, w) := e[i];
10:             newDist := dist + w
11:             **if** (newDist < dis[v]) **then**
12:                 dis[v] := newDist
13:                 PQ.insert(newDist, u)

applied when the graph is sparse.

Recall that the priority-queue-based Dijkstra's algorithm has to update the weight whenever a shorter path is found to any vertex. In an oblivious version of Dijkstra's, this operation dominates the overhead, as it is unclear how to realize it more efficiently than using generic ORAMs. Our solution to oblivious SSSP is more efficient thanks to (1) avoiding this

---

Algorithm 2 Oblivious Dijkstra' algorithm

Secret Input: e, s: same as Algorithm 1
Output: dis: the shortest distance from s to each node
1: dis := $[\infty, \infty, ..., \infty]$; dis[source] = 0
2: PQ.push(0, s); innerLoop := false
3: **for** $i := 0 \to 2V + E$ **do**
4:     **if** not innerLoop **then**
5:         (dist, u) := PQ.deleteMin()
6:         **if** dis[u] == dist **then**
7:             dis[u] := −dis[u]; i := s[u]
8:             innerloop := true;
9:     **else**
10:         **if** i < s[u + 1] **then**
11:             (u, v, w):= e[i]
12:             newDist := dist + w
13:             **if** newDist < dis[u] **then**
14:                 dis[u] := newDist
15:                 PQ.insert(newDist, v')
16:             i = i + 1
17:         **else**
18:             innerloop := false;

weight update operation, and (2) a *loop coalescing* technique.

Avoiding weights updating. This is accomplished by two changes to a standard priority-queue-based Dijkstra's algorithm, i.e., lines 6-7 and line 12 in Algorithm 1. The basic idea is, whenever a shorter distance newDist from s to a vertex u is found, instead of updating the existing weight of u in the heap, we insert a new pair (newDis, u) into the priority

queue. This change can result in multiple entries for the same vertex in the queue, leading to two concerns: (1) the size of the priority queue cannot be bounded by V; and (2) the same vertex might be popped and processed multiple times from the queue. Regarding the first concern, we note the size of the queue can be bounded by $E = O(V^2)$ (since $E = o(V^2)$ for sparse graphs). Hence, each priority queue `insert` and `deleteMin` operation can still be implemented obliviously in $O(\log^2 V)$ [26]. The second concern is resolved by the check in lines 6-7: every vertex will be processed at most once because `dis[v]` will be set negative once vertex v is processed.

Loop coalescing. In Algorithm 1, the two nested loops (line 4 and line 8) use secret data as guards. In order not to leak the secret loop guards, a naive approach is to iterate each loop a maximal number of times (i.e., $V+E$, as V alone is considered secret).

Using our loop coalescing technique, we can derive an oblivious Dijkstra's algorithm that asymptotically outperforms a generic ORAM baseline for sparse graphs. The resulting oblivious algorithm is described in Algorithm 2. Note that at most V vertices and E edges will be visited, we coalesce the two loops into a single one. The code uses a state variable `innerloop` to indicate whether a vertex or an edge is being processed. Each iteration deals with one of a vertex (lines 5-8), an edge (lines 15-18), or the end of a vertex's edges (line 13). So there are $2V + E$ iterations in total. Note the ObliVM-lang compiler will pad the `if`-branches in Algorithm 2 to ensure obliviousness. Further, an oblivious priority queue is employed for `PQ`.

Cost analysis. In Algorithm 2, each iteration of the loop (lines 3-18) makes a constant number of ORAM accesses and two priority queue primitives (`insert` and `deleteMin`, both implemented in $O(\log^2 V)$ time). So, the total runtime is $O((V + E)\log^2 V)$.

Additional algorithmic results. Summarized in Table II, our loop coalescing technique also immediately gives a new oblivious Minimum Spanning Tree (MST) algorithm whose full description is deferred to Appendix A. Additionally, in the process of developing rich libraries for ObliVM, we also designed a novel oblivious Depth First Search (DFS) algorithm that asymptotically outperforms a generic ORAM baseline for dense graphs. The new DFS requires new algorithmic techniques, and we defer its full description to Appendix A.

V. IMPLEMENTING RICH CIRCUIT LIBRARIES IN SOURCE LANGUAGE

A. Case Study: Basic Arithmetic Operations

The rich language features provided by ObliVM-lang make it possible to implement complex arithmetic operations easily and efficiently. We give a case study to demonstrate how to use ObliVM-lang to implement Karatsuba multiplication.

Implementing Karatsuba multiplication. Figure 4 contains the implementation of Karatsuba multiplication [50] in ObliVM-lang. Karatsuba multiplication implements the following recursive algorithm to compute multiplication of two n bit numbers, x and y, taking $O(n^{\log_2 3})$ amount of time. As a quick overview, the algorithm works as follows. First,

```
1   int@(2*n) karatsubaMult@n(
      int@n x, int@n y) {
2     int@2*n ret;
3     if (n < 18) {
4       ret = x*y;
5     } else {
6       int@(n-n/2) a = x$n/2~n$;
7       int@(n/2) b = x$0~n/2$;
8       int@(n-n/2) c = y$n/2~n$;
9       int@(n/2) d = y$0~n/2$;
10      int@(2*(n-n/2)) t1 =
          karatsubaMult@(n-n/2)(a, c);
11      int@(2*(n/2)) t2 =
          karatsubaMult@(n/2)(b, d);
12      int@(n-n/2+1) aPb = a + b;
13      int@(n-n/2+1) cPd = c + d;
14      int@(2*(n-n/2+1)) t3 =
          karatsubaMult@(n-n/2+1)(aPb, cPd);
15      int@(2*n) padt1 = t1;
16      int@(2*n) padt2 = t2;
17      int@(2*n) padt3 = t3;
18      ret = (padt1<<(n/2*2)) + padt2 +
          ((padt3 - padt1 - padt2)<<(n/2));
19    }
20    return ret;
21  }
```

Fig. 4: **Karatsuba multiplication in ObliVM-lang.** In line 6 to line 9, x$i~j$ is used to extract the i-th to the j-th bits of x.

express the n-bit integers x and y as the concatenation of $\frac{n}{2}$-bit integers: x = a*2^{n/2}+b, y = c*2^{n/2}+d. Now, x*y can be calculated as follows:

```
t1 = a*c; t2 = b*d; t3 = (a+b)*(c+d);
x*y = t1<<n + t2 + (t3-t1-t2)<<(n/2);
```

where the multiplications a*c and b*d are implemented through a recursive call to the Karatsuba algorithm itself (until the bit-length is small enough).

To implement Karatsuba efficiently, we need to perform operations on a subset of bits. We hence introduce the following syntactic sugar in ObliVM-lang: In lines 6 to 9 of Figure 4, the syntax num$i~j$ means extracting the part of integer num from i-th bit to j-th bit.

B. Case Study: Circuit ORAM

In Figure 5, we show part of the Circuit ORAM implementation using ObliVM-lang. Line 3 to line 6 is the definition of a ORAM block containing two metadata fields of an index of type int, and a position label of type `rnd`, along with a data field of type `<T>`.

Circuit ORAM (line 7-10) is organized to contain an array of buckets (i.e. arrays of ORAM blocks), and a stash (i.e. an array of blocks). The `dummy` keyword in front of `Block@n<T>` indicates the value of this type can be `null`. In many cases, (e.g. Circuit ORAM implementation), using `dummy` keyword leads to a more efficient circuit generation.

Line 11-30 demonstrates how `readAndRemove` can be implemented. Taking the input of an secret integer index id, and a random position label pos, the label pos is first declassified

```
1   #define BUCSIZE 3
2   #define STASHSIZE 33
3   struct Block@n<T>{
4     int@n id, pos;
5     T   data;
6   };
7   struct CircuitOram@n<T>{
8     dummy Block@n<T>[public 1<<n+1]
9        [public BUCSIZE] buckets;
9     dummy Block@n<T>[public STASHSIZE] stash;
10  };
11  phantom T  CircuitOram@n<T>
        .readAndRemove(int@n id, rnd@n pos) {
12    public int32 pubPos = pos;
13    public int32 i = (1 << n) + pubPos;
14    T  res;
15    for (public int32 k = n; k>=0; k=k-1) {
16      for (public int32 j=0;j<BUCSIZE;j=j+1)
17        if (buckets[i][j] != null &&
18             buckets[i][j].id == id){
19           res = buckets[i][j].data;
20           buckets[i][j] = null;
21        }
22      i=(i-1)/2;
23    }
24    for (public int32 i=0;i<STASHSIZE;i=i+1)
25      if (stash[i]!=null&&stash[i].id==id) {
26         res = stash[i].data;
27         stash[i] = null;
28      }
29    return res;
30  }
```

Fig. 5: **Part of our Circuit ORAM implementation in ObliVM-lang.** In Line 12, pos with type rnd is assigned to a public integer. Here, an implicit declassification on pos happens. The affine type system enforces that pos can be implicitly declassified at most once.

into public. Then affine type system allows declassifying pos once, i.e. pos is never used for the rest of the program. Further in a function calling readAndRemove with inputs arg1 and arg2, arg2 cannot be used either for the rest of the program. This is crucial to enforce that every position labels will use revealed only once after its generation, and, to our best knowledge, no prior work enables such an enforcement in a compiler.

VI. Back End Architecture

Our compiler emits code to a Java-based secure computation back end called ObliVM-GC. We defer details of ObliVM-GC to our online full version [39].

VII. Evaluation

A. Metrics and Experiment Setup

Number of AND gates. In Garbled Circuit-based secure computation, functions are represented in boolean circuits consisting of XOR and AND gates. Due to well-known Free XOR techniques [5]–[7], the cost of evaluating XOR gates are insignificant in comparison with AND gates. Therefore, a primary performance metric is *the number of AND gates*. This metric is *platform independent*, i.e., independent of the artifacts of the underlying software implementation, or the hardware configurations where the benchmark numbers are measured. This metric facilitates a fair comparison with existing works based on boolean circuits, and is one of the most popular metrics used in earlier works [10], [11], [15], [16], [25], [26], [33], [51], [52].

Wall-clock runtime. Unless noted otherwise, all wall-clock numbers are measured by executing the protocols between two Amazon EC2 machines of types c4.8xlarge and c3.8xlarge. This metric is platform and implementation dependent, and therefore we will explain how to best interpret wallclock runtimes, and how these runtimes will be affected by the underlying software and hardware configurations.

Compilation time. For all programs we ran, the compilation time is under 1 second. Therefore, we do not separately report the compilation time for each program.

B. Comparison with Previous Automated Approaches

The first general-purpose secure computation system, Fairplay, was built in 2004 [12]. Since then, several improved systems were built [9]–[11], [13], [14], [16], [33]. Except for our prior work SCVM [15], existing systems provide no support for ORAM – and therefore, each dynamic memory access would be compiled to a linear scan of memory.

We now evaluate the speedup ObliVM achieves relative to previous approaches. To illustrate the sources of the speedup, we consider the following sequence of progressive baselines. We start from Baseline 1 which is representative of a state-of-the-art automated secure computation system. We then add one feature at a time to the baseline, resulting in the next baseline, until we arrive at Baseline 5 which is essentially our ObliVM system.

- **Baseline 1: A state-of-the-art automated system with no ORAM support.** Baseline 1 is intended to characterize a state-of-the-art automated secure computation system with no ORAM support. We assume a compiler that can detect public memory accesses (whose addresses are statically inferrable), and directly make such memory accesses. For each each dynamic memory access (whose address depends on secret inputs), a linear scan of memory is employed. Baseline 1 is effectively a lower-bound estimate of the cost incurred by CMBC-GC [16], a state-of-the-art system in 2012.
- **Baseline 2: With GO-ORAM [40].** In Baseline 2, we implement the GO-ORAM scheme on top of Baseline 1. Dynamic memory accesses made by a program will be compiled to GO-ORAM accesses. We make no additional compile-time optimizations.
- **Baseline 3: With Circuit ORAM [29].** Baseline 3 is essentially the same as Baseline 2 except that we now replace the ORAM scheme with a state-of-the-art Circuit ORAM scheme [29].
- **Baseline 4: Language and compiler.** Baseline 4 assumes that the ObliVM language and compiler are additionally employed (on top of Baseline 3), resulting in additional savings due to our compile-time optimizations as well as our oblivious programming abstractions.

	Oblivious programming abstractions and compiler optimizations demonstrated	Parameters for Figure 6	Parameters for Table IV and Table V
Dijkstra's Algorithm MST	Loop coalescing abstraction (see Section IV-C).	$V = 2^{14}, E = 3V$	$V = 2^{10}, E = 3V$
Heap Map/Set Binary Search	Oblivious data structure abstraction (see Section IV-B).	$N = 2^{27}, K = 32, V = 480$ $N = 2^{23}, K = 32, V = 992$	$N = 2^{23}, K = 32, V = 992$
AMS Sketch	Compile-time optimizations: split data into separate ORAMs [15].	$\epsilon = 6 \times 10^{-5}, \delta = 2^{-20}$	$\epsilon = 2.4 \times 10^{-4}, \delta = 2^{-20}$
Count Min Sketch		$\epsilon = 3 \times 10^{-6}, \delta = 2^{-20}$	
K-Means	MapReduce abstraction (see Section IV-A).	$N = 2^{18}$	$N = 2^{16}$

TABLE III: **List of applications used in Figures 6.** For graph algorithms, V, E stand for number of vertices and edges; for data structures, N, K, V stand for capacity, bit-length of key and bit-length of value; for streaming algorithms, ϵ, δ stand for relative error and failure probability; for K-Means, N stands for number of points.

- **Baseline 5: Back end optimizations.** In Baseline 5, we employ additional back end optimizations atop Baseline 4. Baseline 5 reflects the performance of the actual ObliVM system.

We consider a set of applications in our evaluation as described in Table III. We select several applications to showcase our oblivious programming abstractions, including MapReduce, loop coalescing, and oblivious data structure abstractions. For all applications, we choose moderately large data sizes ranging from 768KB to 10GB. For data structures (e.g., Heap, Map/Set) and binary search, for Baseline 1, we assume that each operation (e.g., search, add, delete) is done with a single linear scan. For Baseline 2 and 3, we assume that a typical sub-linear implementation is adopted. For all other applications, we assume that Baseline 1 to 3 adopt the most straightforward implementation of the algorithm.

Results. Figure 6 shows the speedup we achieve relative to a state-of-the-art automated system that does not employ ORAM [16]. This speedup comes from the following sources:

No ORAM to GO-ORAM: For most of the cases, the data size considered was not big enough for GO-ORAM to be competitive to a linear-scan ORAM. The only exception was AMS sketch, where we chose a large sketch size. In this case, using GO-ORAM would result in a $300\times$ speedup in comparison with no ORAM (i.e., linear-scan for each dynamic memory access). This part of the speedup is reflected in purple in Figure 6. Here the speedup stems from a reduction in circuit size (as measured by the number of AND gates).

Circuit ORAM: The red parts in Figure 6 reflect the multiplicative speedup attained when we instead use Circuit ORAM (as opposed to no ORAM or GO-ORAM, whichever is faster). This way, we achieve an additional $51\times$ to 530 performance gains – reflected by a reduction in the total circuit size.

Language and compiler: As reflected by the blue bars in Figure 6, our oblivious programming abstractions and compile-time optimizations bring an additional $2\times$ to $2500\times$ performance savings on top of a generic Circuit ORAM-based approach. This speedup is also measurable in terms of reduction in the circuit size.

Back end optimizations: Our ObliVM-GC is a better architected and more optimized version of its predeces-

sor FastGC [33] which is employed by CMBC-GC [16]. FastGC [33] reported a garbling speed of 96K AND gates/sec, whereas ObliVM garbles at 670K AND gates/sec on a comparable machine. In total, we achieve an $7\times$ overall speedup compared with FastGC [33].

We stress, however, that ObliVM's main contribution is not the back end implementation. In fact, it would be faster to hook up ObliVM's language and compiler with a JustGarble-like system that employs a C-based implementation and hardware AES-NI. However, presently JustGarble does not provide a fully working end-to-end protocol. Therefore, it is an important direction of future work to extend JustGarble to a fully working protocol, and integrate it into ObliVM.

Comparison with SCVM. In comparison with SCVM, ObliVM's offers the following new features: 1) new oblivious programming abstractions; 2) Circuit ORAM implementation that is $20\times$ to $30\times$ times faster than SCVM's binary-tree ORAM implementation for 4MB to 4GB data sizes; and 3) ability to implement low-level gadgets including the ORAM algorithm itself in the source language.

In the online full version [39], we give a detailed comparison with an SCVM-like system. Since the design of efficient ORAM algorithms is mainly the contribution of the Circuit ORAM paper [29], here we focus on evaluating the gains from programming abstractions. Therefore, instead of comparing with SCVM per se, we compare with SCVM + Circuit ORAM instead (i.e., SCVM with its ORAM implementation updated to the latest Circuit ORAM).

C. *ObliVM vs. Hand-Crafted Solutions*

We show that ObliVM achieves competitive performance relative to hand-crafted solutions for a wide class of common tasks. We also show that ObliVM significantly reduces development effort in comparison with previous secure computation frameworks.

Competitive performance. For a set of applications, including Heap, Map/Set, AMS Sketch, Count-Min Sketch, and K-Means, we compared implementations auto-generated by ObliVM with implementations hand-crafted by human experts. Here the human experts are authors of this paper. We assume that the human experts have wisdom of employing the most efficient, state-of-the-art oblivious algorithms when designing circuits for these algorithms. For example, Histogram and K-Means algorithms are implemented with oblivious sorting pro-

Fig. 6: **Sources of speedup in comparison with state-of-the-art in 2012 [16]: an in-depth look.**

tools instead of generic ORAM. Heap and Map/Set employ state-of-the-art oblivious data structure techniques [26]. The graph algorithms including Dijkstra and MST employ novel oblivious algorithms proposed in this paper. In comparison, our ObliVM programs for the same applications do not require special security expertise to create. The programmer simply has to express these tasks in the programming abstractions we offer whenever possible. Over the suite of application benchmarks we consider, our ObliVM programs are competitive to hand-crafted implementations – and the performance difference is only **0.5% − 2%** throughout.

We remark that the hand-crafted circuits are not necessarily the optimal circuits for each computation task. However, they do represent asymptotically the best known algorithms (or new algorithms that are direct implications of this paper). It is conceivable that circuit optimization techniques such as those proposed in TinyGarble [34] can further reduce circuit sizes by a small constant factor (e.g., 50%). We leave this part as an interesting direction of future research.

Developer effort. We use two concrete case studies to demonstrate the significant reduction of developer effort enabled by ObliVM.

Case study: ridge regression. Ridge regression [53] takes as input a large number of data points and finds the best-fit linear curve for these points. The algorithm is an important building block in various machine-learning tasks [52]. Previously, Nikolaenko et al. [52] developed a system to securely evaluate ridge regression, using the FastGC framework [33], which took them roughly **three weeks** [54]. In contrast, we spent **two student·hours** to accomplish the same task using ObliVM. In addition to the speedup gain from ObliVM-GC back end, our optimized libraries result in $3\times$ smaller circuits with aligned parameters. We defer the detailed comparison to the online technical report [39].

Case study: oblivious data structures. Oblivious AVL tree (i.e, the Map/Set data structure) is an example algorithm that was previously too complex to program as circuits, but now becomes very easy with ObliVM. In an earlier work [26], we

designed an oblivious AVL tree algorithm, but were unable to implement it due to high programming complexity. Now, with ObliVM, we implement an AVL tree with 311 lines of code in ObliVM-lang, consuming under 10 student·hours (including the implementation as well as debugging).

We stress that it is not possible to implement oblivious AVL tree in previous languages for secure computation, including the state-of-the-art Wysteria [17].

D. End-to-End Application Performance

Currently in ObliVM-GC, we implemented a standard garbling scheme with Garbled Row Reduction [36] and FreeeXOR [5]. We also implemented an OT extension protocol proposed by Ishai et al. [3] and a basic OT protocol by Naor and Pinkas [55].

Setup. For evaluation, here we consider a scenario where a client secret shares its data between two non-colluding cloud providers a priori. For cases where inputs are a large dataset (e.g., Heap, Map/Set, etc), depending on the application, the client may sometimes need to place the inputs in an ORAM, and secret-share the resulting ORAM among the two cloud providers. We do not measure this setup cost in our evaluation – this cost can depend highly on the available bandwidth between the client and the two cloud providers. Therefore, our evaluation begins assuming this one-time setup has completed.

End-to-end application performance. In Table IV, we consider three types of applications, basic instructions (e.g., addition, multiplication, and floating point operations); linear or super-linear algorithms (e.g., Dijkstra, K-Means, Minimum Spanning Tree, and Histogram); and sublinear-time algorithms (e.g., Heap, Map/Set, Binary Search, Count Min Sketch, AMS Sketch). We report the circuit size, online and total costs for a variety of applications at typical data sizes.

In Table IV, we also compare ObliVM with a state-of-the-art automated secure computation system CMBC-GC [16]. We note that the authors of CMBC-GC did not run all of these application benchmarks, so we project the performance of CMBC-GC using the following estimate: we first change

Program	Input size	CMBC-GC (estimate)		ObliVM Framework			ObliVM + JustGarble (estimate)	
		#AND gates	Total time	#AND gates	Total time	Online time	Total time	Online time
Basic instructions								
Integer addition	1024 bits	2977	31ms	**1024**	**1.7ms**	**0.6ms**	0.12ms	0.05ms
Integer mult.	1024 bits	6.4M	66.4s	**572K**	**833ms**	**274ms**	69.4ms	28.9ms
Integer Comparison	16384 bits	32K	335.7ms	**16384**	**26ms**	**8.58ms**	1.96ms	0.82ms
Floating point addition	64 bits	10K	104ms	**3035**	**4.32ms**	**1.45ms**	0.36ms	0.15ms
Floating point mult.	64 bits	10K	104ms	**4312**	**6.29ms**	**2.02ms**	0.52ms	0.22ms
Hamming distance	1600 bits	30K	310ms	**3200**	**5.07ms**	**1.71ms**	0.39ms	0.16ms
Linear or super-linear algorithms								
K-Means	0.5MB	550B	66d	**2269M**	**62.1min**	**23.6min**	4.58min	1.9min
Dijkstra's Algorithm	48KB	755B	91d	**10B**	**12.6h**	**3.09h**	20.4min	8.5min
MST	48KB	755B	91d	**9.6B**	**12.4h**	**3h**	19.6min	8.2min
Histogram	0.25MB	137B	16.5d	**866M**	**21.5min**	**8.56min**	1.7min	42.5s
Sublinear-time algorithms								
Heap	1GB	32B	3.9d	**12.5M**	**59.3s**	**10.42s**	1.5s	625ms
Map/Set	1GB	32B	3.9d	**23.9M**	**117.2s**	**20.67s**	2.9s	1.2s
Binary Search	1GB	32B	3.9d	**1562K**	**7.36s**	**1.34s**	189ms	78.8ms
Count Min Sketch	0.31GB	9.9B	30.8h	**8088K**	**20.77s**	**6.4s**	0.98s	0.41s
AMS Sketch	1.25GB	40B	5.18d	**9949K**	**36.76s**	**9.95s**	1.21s	504ms

TABLE IV: **Application performance.** Actual measured numbers are in bold. The remainder are estimated numbers and should be interpreted with care. ObliVM numbers for basic instructions and sublinear-time algorithms are the mean of 20 runs. Since for all these applications, our measurements have small spread (all runs are within 6% from the mean), we use a single run for linear-time and super-linear algorithms (the same for Table V).

our compiler to adopt a linear scan of memory upon dynamic memory accesses – this allows us to obtain an estimate of the circuit size CMBC-GC would have obtained for the same applications. For the set of application benchmarks (e.g., K-Means, MST, etc) CMBC-GC did report in their paper, we confirmed that our circuit size estimates are always a lower bound of what CMBC-GC reported. We then estimate the runtime of CMBC-GC based on their reported 96K AND gates per sec – assuming that a network bandwidth of at least 2.8MBps is provisioned.

As mentioned earlier, the focus of this paper is our language and compiler, not the back end cryptographic implementation. It should be relatively easy to integrate our language and compiler with a JustGarble-like back end that employs hardware AES-NI. In Table IV, we also give an estimate of the performance we anticipate if we ran our ObliVM-generated circuits over a JustGarble-like back end. This is calculated using our circuit sizes and the 11M AND gates/sec performance number reported by JustGarble [4].

- *Online cost.* To measure online cost, we assume that all work that is independent of input data is performed offline, including garbling and input-independent OT pre-processing. Our present ObliVM implementation achieves an online speed of 1.8M gates/sec consuming roughly 54MBps network bandwidth.
- *Offline cost.* When no work is deferred to an offline phase, ObliVM achieves a garbling speed of $670K$ gates/sec consuming 19MBps network bandwidth.

Slowdown relative to a non-secure baseline. For completeness, we now describe ObliVM's slowdown in comparison with a non-secure baseline where computation is performed in cleartext. As shown in Table V, our slowdown relative to a non-secure baseline is application dependent, and ranges from $45\times$ to $9.3 \times 10^6\times$. We also present the *estimated* slowdown if a JustGarble-like back end is used for ObliVM-generated circuits. These numbers are estimated based on our circuit sizes as well as the reported 11M AND gates/sec performance metric reported by JustGarble [4].

In particular, we elaborate on the following interesting cases. First, the distributed genome-wide association study(GWAS) application is Task 1 in the iDash secure genomic analysis competition [56], with total data size 380KB. This task achieves a small slowdown, because part of the computation is done locally – specifically, Alice and Bob each performs some local preprocessing to obtain the alle frequencies of their own data, before engaging in a secure computation protocol to compute χ^2-statistics. For details, we refer the reader to our online short note on how we implemented the competition tasks. On the other hand, benchmarks with floating point operations such as K-Means incur a relatively larger slowdown because modern processors have special floating point instructions which makes it favorable to the insecure baseline.

VIII. CONCLUSION, SUBSEQUENT AND FUTURE WORK

We design ObliVM, a programming framework for automated secure computation. Additional examples can be found at our project website http://www.oblivm.com, including popular streaming algorithms, graph algorithms, data structures, machine learning algorithms, secure genome analysis [56], etc.

Task Cleartext Time	K-Means 0.4ms		Distributed GWAS 40ms		Binary Search 10μs		AMS Sketch 80μs		Hamming dist. 0.3μs	
	Online	Total	Online	Total	Online	Total	Online	Total	Online	Total
ObliVM Runtime	24min	62min	1.8s	5.2s	1.3s	7.4s	9.5s	36.8s	1.71ms	5.07ms
Slowdown	3.6×10^6	9.3×10^6	45	130	1.3×10^5	7.4×10^5	1.2×10^5	4.6×10^5	6×10^3	1.7×10^4
ObliVM+JustGB (estimate)	1.9min	4.58min	0.14s	0.28s	78.8ms	189ms	0.5s	1.2s	0.16ms	0.39ms
Slowdown	2.9×10^5	6.9×10^5	3.5	7	7.9×10^3	1.9×10^4	6.3×10^3	1.5×10^4	5.3×10^2	1.3×10^3

TABLE V: **Slowdown of secure computation compared with non-secure, cleartext computation.** Parameter choices are the same as Table IV. Online cost only includes operations that are input-dependent. All time measurements assume data are pre-loaded to the memory. ObliVM requires a bandwidth of 19MBps. Numbers for JustGarble are estimated using ObliVM-generated circuit sizes assuming 315MBps bandwidth.

A. Subsequent Works and Adoption of ObliVM

To the best of our knowledge, our framework is already being adopted in other projects. First, the GraphSC work [24] extends our ObliVM-GC framework to support parallel execution of gadgets on modern architectures with inherent parallelism, such as multi-core processor architectures, and compute clusters. Because of ObliVM-GC's clean architecture, it was not too much work for GraphSC's parallel extension, which required about 1200 more lines of code on top of ObliVM-GC. Second, our collaborators (and a subset of the authors of this paper) are implementing a MIPS processor over our ObliVM framework. Such a MIPS processor will allow maximum backward compatibility: code written in any language can be compiled to a MIPS processor using a stock compiler, and then evaluated securely. Third, a group of networking researchers have used our ObliVM-GC framework to develop privacy-preserving software-defined networking applications [57]. Fourth, we used our ObliVM framework to participate in the iDash Secure Genome Analysis Competition [56], [58]. Finally, Wagner et al. also use our ObliVM framework to develop privacy-preserving applications on human microbiomes [59].

B. Future Work

In future work, we will implement a C-based Garbled Circuit back end similar to JustGarble [4], such that we can exploit hardware AES-NI features of modern processors. We will also implement the state-of-the-art OT optimizations [8]. It will also be interesting to provide support for multiple parties and malicious security. Since ObliVM is designed to be good at compiling programs to compact circuits, it will be interesting to extend ObliVM to support other cryptographic back ends such as fully homomorphic encryption, program obfuscation and verifiable computation.

ACKNOWLEDGMENTS

REFERENCES

[1] A. C.-C. Yao, "Protocols for secure computations (extended abstract)," in *FOCS*, 1982.

[2] O. Goldreich, S. Micali, and A. Wigderson, "How to play any mental game," in *STOC*, 1987.

[3] Y. Ishai, J. Kilian, K. Nissim, and E. Petrank, "Extending Oblivious Transfers Efficiently," in *CRYPTO 2003*, 2003.

[4] M. Bellare, V. T. Hoang, S. Keelveedhi, and P. Rogaway, "Efficient Garbling from a Fixed-Key Blockcipher," in *S & P*, 2013.

[5] V. Kolesnikov and T. Schneider, "Improved Garbled Circuit: Free XOR Gates and Applications," in *ICALP*, 2008.

[6] S. G. Choi, J. Katz, R. Kumaresan, and H.-S. Zhou, "On the security of the "free-xor" technique," in *TCC*, 2012.

[7] B. Applebaum, "Garbling xor gates "for free" in the standard model," in *TCC*, 2013.

[8] G. Asharov, Y. Lindell, T. Schneider, and M. Zohner, "More Efficient Oblivious Transfer and Extensions for Faster Secure Computation," ser. CCS '13, 2013.

[9] D. Bogdanov, S. Laur, and J. Willemson, "Sharemind: A Framework for Fast Privacy-Preserving Computations," in *ESORICS*, 2008.

[10] B. Kreuter, B. Mood, A. Shelat, and K. Butler, "PCF: A portable circuit format for scalable two-party secure computation," in *Usenix Security*, 2013.

[11] B. Kreuter, a. shelat, and C.-H. Shen, "Billion-gate secure computation with malicious adversaries," in *USENIX Security*, 2012.

[12] D. Malkhi, N. Nisan, B. Pinkas, and Y. Sella, "Fairplay: a secure two-party computation system," in *USENIX Security*, 2004.

[13] W. Henecka, S. Kögl, A.-R. Sadeghi, T. Schneider, and I. Wehrenberg, "Tasty: tool for automating secure two-party computations," in *CCS*, 2010.

[14] Y. Zhang, A. Steele, and M. Blanton, "PICCO: a general-purpose compiler for private distributed computation," in *CCS*, 2013.

[15] C. Liu, Y. Huang, E. Shi, J. Katz, and M. Hicks, "Automating Efficient RAM-model Secure Computation," in *S & P*, May 2014.

[16] A. Holzer, M. Franz, S. Katzenbeisser, and H. Veith, "Secure Two-party Computations in ANSI C," in *CCS*, 2012.

[17] A. Rastogi, M. A. Hammer, and M. Hicks, "Wysteria: A Programming Language for Generic, Mixed-Mode Multiparty Computations," in *S & P*, 2014.

[18] "Partisia," http://www.partisia.dk/.

[19] "Dyadic security," http://www.dyadicsec.com/.

[20] R. Canetti, "Security and composition of multiparty cryptographic protocols," *Journal of Cryptology*, 2000.

[21] C. Liu, X. S. Wang, M. Hicks, and E. Shi, "Formalizing the ObliVM language," Manuscript in preparation, 2015.

[22] M. T. Goodrich and J. A. Simons, "Data-Oblivious Graph Algorithms in Outsourced External Memory," *CoRR*, vol. abs/1409.0597, 2014.

[23] J. Brickell and V. Shmatikov, "Privacy-preserving graph algorithms in the semi-honest model," in *ASIACRYPT*, 2005.

[24] K. Nayak, X. S. Wang, S. Ioannidis, U. Weinsberg, N. Taft, and E. Shi, "GraphSC: Parallel Secure Computation Made Easy," in *IEEE S & P*, 2015.

[25] V. Nikolaenko, S. Ioannidis, U. Weinsberg, M. Joye, N. Taft, and D. Boneh, "Privacy-preserving matrix factorization," in *CCS*, 2013.

[26] X. S. Wang, K. Nayak, C. Liu, T.-H. H. Chan, E. Shi, E. Stefanov, and Y. Huang, "Oblivious Data Structures," in *CCS*, 2014.

[27] M. Keller and P. Scholl, "Efficient, oblivious data structures for MPC," in *Asiacrypt*, 2014.

[28] J. C. Mitchell and J. Zimmerman, "Data-Oblivious Data Structures," in *STACS*, 2014, pp. 554–565.

[29] X. S. Wang, T.-H. H. Chan, and E. Shi, "Circuit ORAM: On Tightness of the Goldreich-Ostrovsky Lower Bound," Cryptology ePrint Archive, Report 2014/672, 2014.

[30] "Rsa distributed credential protection," http://www.emc.com/security/rsa-distributed-credential-protection.htm.

[31] C. Liu, M. Hicks, and E. Shi, "Memory trace oblivious program execution," ser. CSF '13, 2013, pp. 51–65.

[32] S. D. Gordon, A. McIntosh, J. Katz, E. Shi, and X. S. Wang, "Secure computation of MIPS machine code," Manuscript, 2015.

[33] Y. Huang, D. Evans, J. Katz, and L. Malka, "Faster secure two-party computation using garbled circuits," in *Usenix Security Symposium*, 2011.

[34] E. M. Songhori, S. U. Hussain, A.-R. Sadeghi, T. Schneider, and F. Koushanfar, "TinyGarble: Highly Compressed and Scalable Sequential Garbled Circuits," in *IEEE S & P*, 2015.

[35] A. C.-C. Yao, "How to generate and exchange secrets," in *FOCS*, 1986.

[36] M. Naor, B. Pinkas, and R. Sumner, "Privacy preserving auctions and mechanism design," ser. EC '99, 1999.

[37] S. Zahur, M. Rosulek, and D. Evans, "Two halves make a whole: Reducing data transfer in garbled circuits using half gates," in *EUROCRYPT*, 2015.

[38] C. Gentry, "Fully homomorphic encryption using ideal lattices," in *STOC*, 2009.

[39] C. Liu, X. S. Wang, K. Nayak, Y. Huang, and E. Shi, "Oblivm: A programming framework for secure computation," 2015, http://www.cs.umd.edu/~elaine/docs/oblivmtr.pdf.

[40] O. Goldreich and R. Ostrovsky, "Software protection and simulation on oblivious RAMs," *J. ACM*, 1996.

[41] O. Goldreich, "Towards a theory of software protection and simulation by oblivious RAMs," in *STOC*, 1987.

[42] E. Shi, T.-H. H. Chan, E. Stefanov, and M. Li, "Oblivious RAM with $O((\log N)^3)$ worst-case cost," in *ASIACRYPT*, 2011.

[43] J. Agat, "Transforming out timing leaks," in *POPL*, 2000.

[44] A. Russo, J. Hughes, D. A. Naumann, and A. Sabelfeld, "Closing internal timing channels by transformation," in *ASIAN*, 2006.

[45] M. T. Goodrich and M. Mitzenmacher, "Privacy-preserving access of outsourced data via oblivious RAM simulation," in *ICALP*, 2011.

[46] D. Dachman-Soled, C. Liu, C. Papamanthou, E. Shi, and U. Vishkin, "Oblivious network RAM," Cryptology ePrint Archive, Report 2015/073, 2015, http://eprint.iacr.org/.

[47] G. Malewicz, M. H. Austern, A. J. Bik, J. C. Dehnert, I. Horn, N. Leiser, and G. Czajkowski, "Pregel: a system for large-scale graph processing," in *SIGMOD*, 2010.

[48] "Graphlab," http://graphlab.org.

[49] M. Blanton, A. Steele, and M. Alisagari, "Data-oblivious graph algorithms for secure computation and outsourcing," in *ASIA CCS*, 2013.

[50] A. A. Karatsuba, "The Complexity of Computations," 1995.

[51] X. S. Wang, Y. Huang, T.-H. H. Chan, A. Shelat, and E. Shi, "SCO-RAM: Oblivious RAM for Secure Computation," in *CCS*, 2014.

[52] V. Nikolaenko, U. Weinsberg, S. Ioannidis, M. Joye, D. Boneh, and N. Taft, "Privacy-preserving ridge regression on hundreds of millions of records," in *S & P*, 2013.

[53] T. Hastie, R. Tibshirani, and J. Friedman, *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*, 2001.

[54] "Private communication," 2014.

[55] M. Naor and B. Pinkas, "Efficient oblivious transfer protocols," in *SODA*, 2001.

[56] http://humangenomeprivacy.org/2015.

[57] N. A. Jagadeesan, R. Pal, K. Nadikuditi, Y. Huang, E. Shi, and M. Yu, "A secure computation framework for SDNs," in *HotSDN*, 2014.

[58] X. S. Wang, C. Liu, K. Nayak, Y. Huang, and E. Shi, "idash secure genome analysis competition using oblivm," Cryptology ePrint Archive, Report 2015/191, 2015, http://eprint.iacr.org/.

[59] J. Wagner, J. Paulson, X. S. Wang, H. Corrada-Bravo, and B. Bhattacharjee, "Privacy-preserving human microbiome analysis using secure computation," Manuscript, 2015.

[60] E. Kushilevitz, S. Lu, and R. Ostrovsky, "On the (in)security of hash-based oblivious RAM and a new balancing scheme," in *SODA*, 2012.

APPENDIX A
ADDITIONAL OBLIVIOUS ALGORITHM

A. Additional Oblivious Graph Algorithm

It is has been an open question how to construct an Oblivious Depth First Search (ODFS) algorithm that outperforms one built on generic ORAMs [22]. Here we answer this question for dense graphs. We present $O((E + V) \log V)$ time ODFS algorithm. In comparison, a generic-ORAM based oblivious solution would take $O((E+V) \log^2 V)$ time (ignoring possible $\log \log$ factors) [29], [60].

The challenge is that in standard DFS, we need to verify whether a vertex has been visited every time we explore a new edge. Typically, this is done by storing a bit-array that supports dynamic access. To make it oblivious would require placing this bit-array inside an ORAM, thus incurring $O(\log^2 V)$ cost per access, and $O(E \log^2 V)$ time over all $O(E)$ accesses.

To solve this problem, instead of verifying if a vertex has been visited, we maintain a `tovisit` list of vertexes, which preserves the same traversal order as DFS. When new vertexes are added to this list, we guarantee that each vertex appears in the list at most once using an oblivious sorting algorithm. Algorithm 3 presents our oblivious DFS algorithm, and defines the inputs, outputs, and how they are stored.

Since DFS explores the latest visited vertex first, so we maintain a stack-like `tovisit` array, where the top of the stack is stored in position 1. Each cell of `tovisit` is a pair $(u,$ `depth`$)$:

- $(u,$ `depth` $= \infty)$: indicates that vertex u has been expanded, and will not be expanded again.
- $(u,$ `depth` $\neq \infty)$: indicates that vertex u was reached at depth `depth`. The bigger the `depth`, the sooner u should be expanded.

Each iteration of the main loop (Lines 2-12) reads the top of the stack-like `tovisit` array, and expands the vertex encountered. The most interesting part of the algorithm is Line 12, highlighted in red. In this step, the newly reached vertices in this iteration, stored in the `add` array, will be added to the `tovisit` array in a non-trivial manner as explained below. At

Algorithm 3 Oblivious DFS

Secret Input: s: starting vertex;
Secret Input: E: adjacency matrix, stored in an ORAM of V blocks, *each block being one row of the matrix.*
Output: order: DFS traversal order *// not in ORAM*

```
 1: tovisit:=[(s,0), ⊥, ..., ⊥];  // not in ORAM
 2: for i = 1 → |V| do
 3:     (u, depth) := tovisit[1];
 4:     tovisit[1] := (u, ∞); // mark as visited
 5:     order[i] := u;
 6:     edge := E[u];
 7:     for v := 1 → |V| do
 8:         if edge[v] == 1 then // (u, v) is an edge
 9:             add[v] := (v, i); // add is not in ORAM
10:         else   // (u, v) is not an edge
11:             add[v] := ⊥;
12:     tovisit.Merge(add) ;
13: return order
```
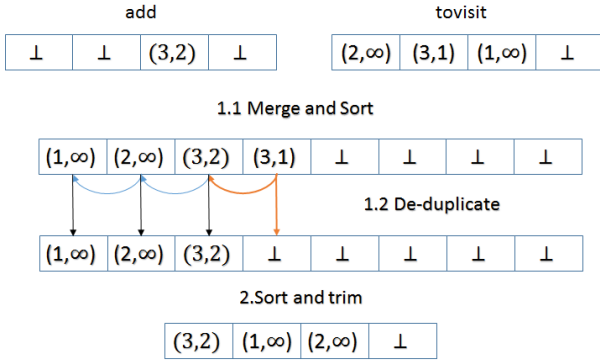

Fig. 7: Oblivious DFS Example: illustration of `tovisit.Merge(add)`.

the end of each iteration (i.e., after executing Line 12), the following invariants hold for the array `tovisit`:

- *Sorted by depth.* All entries in `tovisit` are sorted by their `depth` in decreasing order. This ensures an entry added last (with largest `depth`) will be "popped" first.
- *Visited vertexes will never be expanded.* All entries with a ∞ depth come after those with a finite depth.
- *No duplicates.* Any two entries (v, d) and (v, d') where $(d > d')$ will be combined into (v, d).
- *Fixed length.* The length of `tovisit` is exactly V.

The merge operation (Line 12). The operation is performed with two oblivious sorts. See Figure 7 for an illustrated example.

1) *O-sort and deduplicate.* This sorting groups all entries for the same vertex together, with the `depth` field in descending order (∞ comes first). All $\perp$ entries are moved to the end. Then, for all entries with the same vertex number (which are adjacent), we keep only the first one (which has the largest depth value) while overwriting others with $\perp$.

2) *O-sort and trim.* This sorting will (*a*) push all $\perp$ entries to the end; (*b*) push all ∞ entries to the end; and (*c*) sort all remaining entries in descending order of `depth`. Discard everything but the first V entries.

Algorithm 4 Minimum Spanning Tree with bounded for

Secret Input: s: the source node
Secret Input: e: concatenation of adjacency lists stored in a single ORAM array. Each vertex's neighbors are stored adjacent to each other.
Secret Input: s[u]: sum of out-degree over vertices from 1 to u.
Output: dis: the shortest distance from source to each node

```
 1: explored := [false, false, ..., false]
 2: PQ.push(0, s)
 3: res := 0
 4: bwhile(V)(!PQ.empty())
 5:     (weight, u) := PQ.deleteMin()
 6:     if(!explored[u]) then
 7:         res:= res + weight
 8:         explored[u] := true
 9:         bfor(E)(i := s[u]; i < s[u + 1]; i = i + 1)
10:             (u, v, w) = e[i];
11:             PQ.insert(w, v)
```

Cost analysis. The inner loop (lines 8-11) runs in constant time, and will run V^2 times. Lines 3-5 also run in constant time, but will only run V times. Line 6 is an ORAM read, and it will run V times. Since the ORAM's block size is $V = \omega(\log^2 V)$, each ORAM read has an amortized cost of $O(V \log V)$. Finally, Line 12, which will run V times, consists of four oblivious sortings over an $O(V)$-size array, thus costs $O(V \log V)$. Hence, the overall cost of our algorithm is $O(V^2 \log V)$.

B. Oblivious Minimum Spanning Tree

In Algorithm 4, we show the pseudo-code for minimum spanning tree algorithm written using ObliVM-lang with our new loop coalescing abstraction. The algorithm is very similar to the standard textbook implementation except for the annotations used for bounded-for loops in Lines 4 and 9. As described in Section IV-C, the inner loop (Line 9 to Line 11) will only execute $O(V + E)$ times over all iterations of the outer loop. Further, each execution of the inner loop requires circuits of size $O(\log^2 V)$, using latest oblivious data structures [26] and Circuit ORAM [29]. So the overall complexity is $O((V + E) \log^2 V)$. We defer further discussions about minimum spanning tree algorithm to our online full version [39].